

Safety Fine-tuning: Training Loss (AdvBench + HarmBench refusals)

